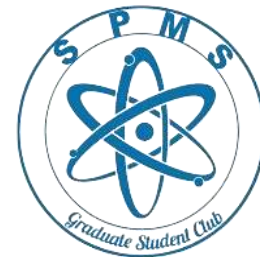


Deep Learning Bootcamp

2026



Agenda

1. Introduction to Python
2. Basic numerical libraries for ML

Break 15:30

1. Linear Regression
2. Exploratory data analysis
3. Linear regression with PyTorch (deep learning library)

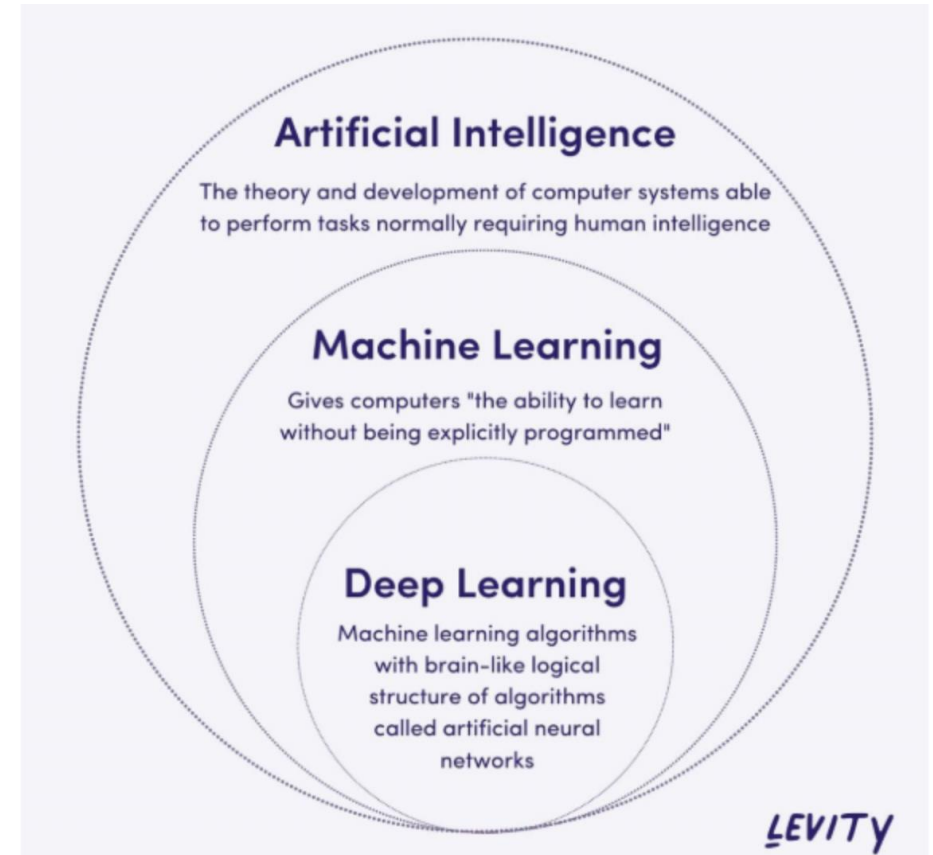
Schedule

Event No.	Details
Event 1 24th Oct 2025	<i>Deep Learning Essentials</i> Covers the basics of Python and necessary packages required for Deep Learning such as numpy, scipy, pandas etc.
Event 2 10th Nov 2025	<i>Deep Learning for Regression and Classification</i> Will cover the basics of PyTorch, as well as how to use PyTorch for performing regression and classification tasks.
Event 3 17th Nov 2025	<i>Deep Learning for Images</i> In this event, we will extend the classification using deep learning, specifically focusing on datasets involving images.
Event 4 TBD – Sem 2 – 2026	<i>Deep Learning for Sequence Data (text and time series)</i> In this event, we will focus on using Deep Learning models for datasets involving sequences or temporal relations. We plan to cover examples from both text and time-series datasets.
Event 5 TBD – Sem 2 – 2026	<i>Reinforcement Learning</i> This session will introduce Deep Reinforcement Learning techniques with some practical applications.

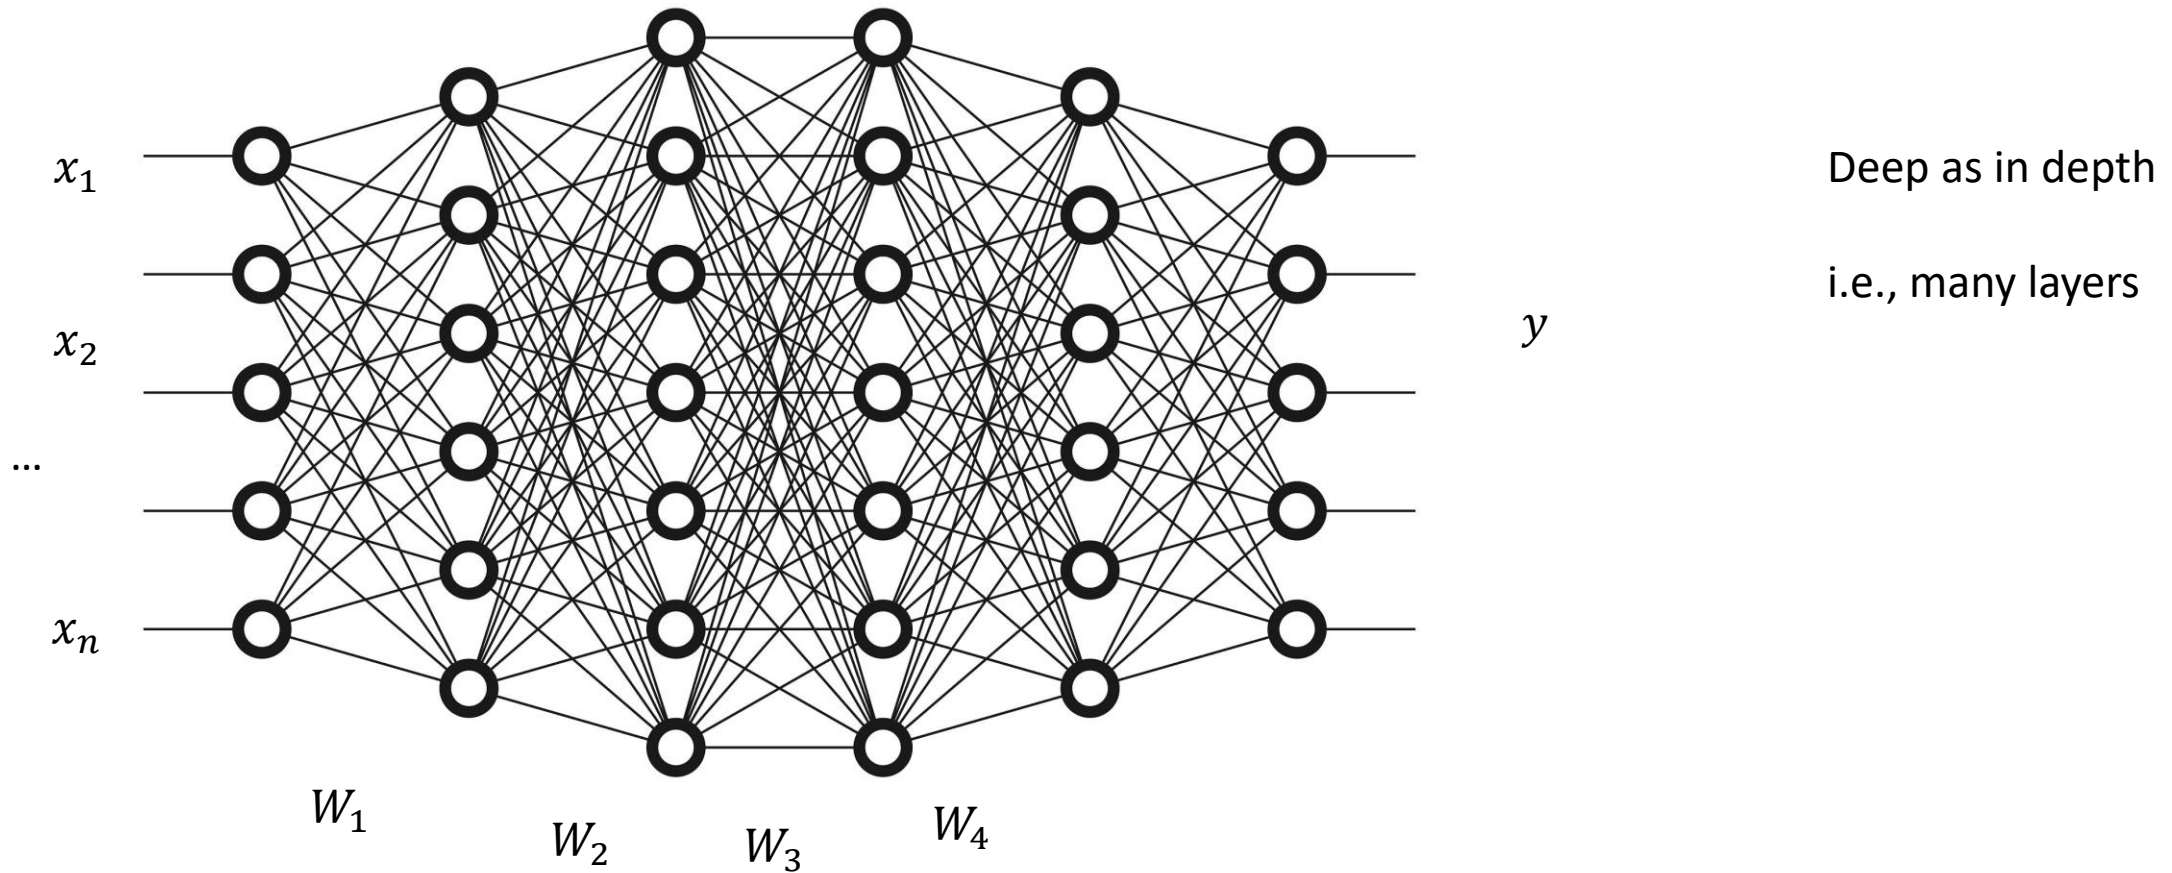
What is Deep Learning?

- Machine learning: learn from examples or data
- Deep Learning: use artificial neural networks: weight * variable + some non-linear layer

$$y = \sigma \left(W_L \dots \sigma(W_2 \sigma(W_1 x)) \right)$$



Deep Learning is powered by deep neural networks



https://cdn-images-1.medium.com/max/2400/1*1mpE6fsq5LNxH31xeTWi5w.jpeg

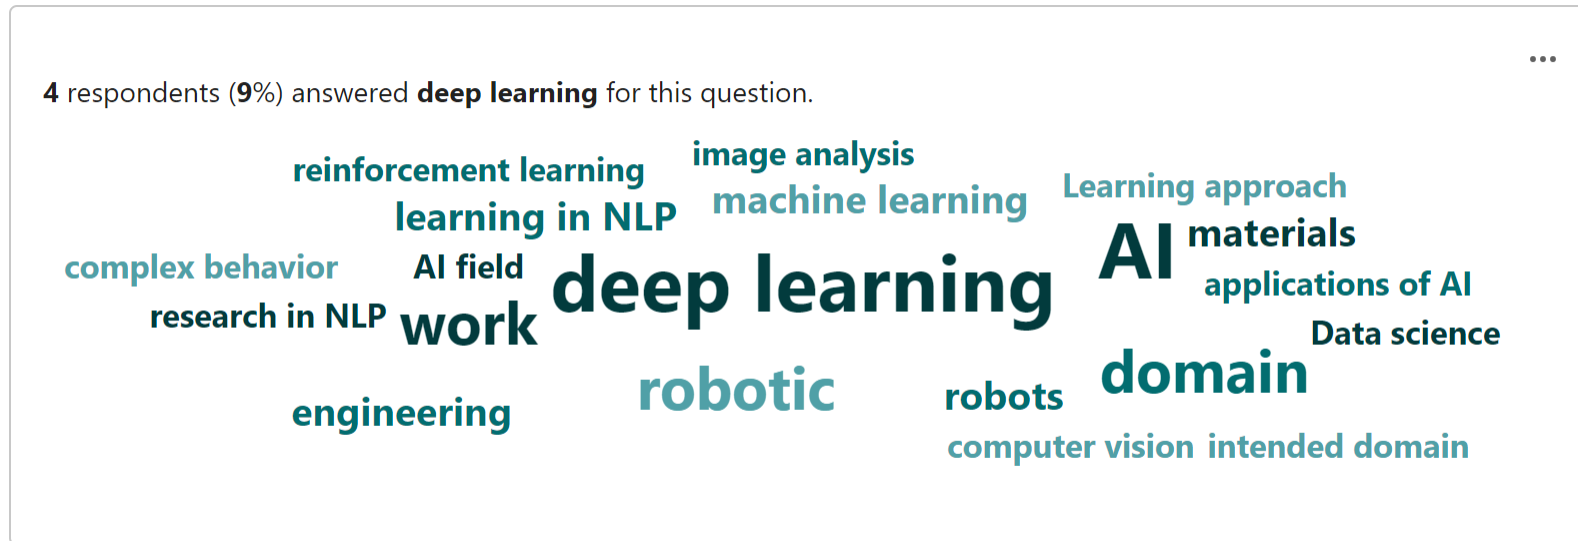
Why Deep Learning?

- Find hidden patterns in a lot of data (high-dimensional, complex, etc.)
- Many applications, open models
 - Material properties prediction
 - Protein folding: see AlphaFold
 - Text generation/Natural Language Processing: ChatGPT
 - Image generation: DALL·E, MidJourney, etc.
 - Computer Vision: ImageNet, image segmentation, detect face in phone, etc.
 - ...

Deep Fakes...



Your interest



Examples

You

suggest a simple dataset to use to predict material property using deep

ChatGPT

For predicting material properties using deep learning, you might consider the Materials Project database, which provides extensive information on various materials. Let's formulate this as a regression problem, where we aim to predict a material property (e.g., band gap) based on certain features.

Dataset:

Materials Project Database - Band gap prediction.

Mathematical Formulation:

Given a dataset with N samples, where each sample is represented by a feature vector X_i and the corresponding material property (band gap) y_i , the goal is to find a mapping function f such that $f(X_i) \approx y_i$.

Python Code (using PyTorch):

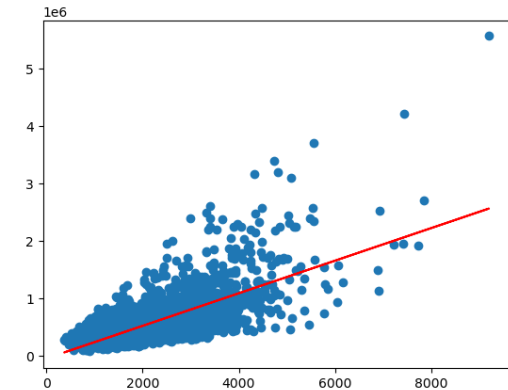
```
python

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
```

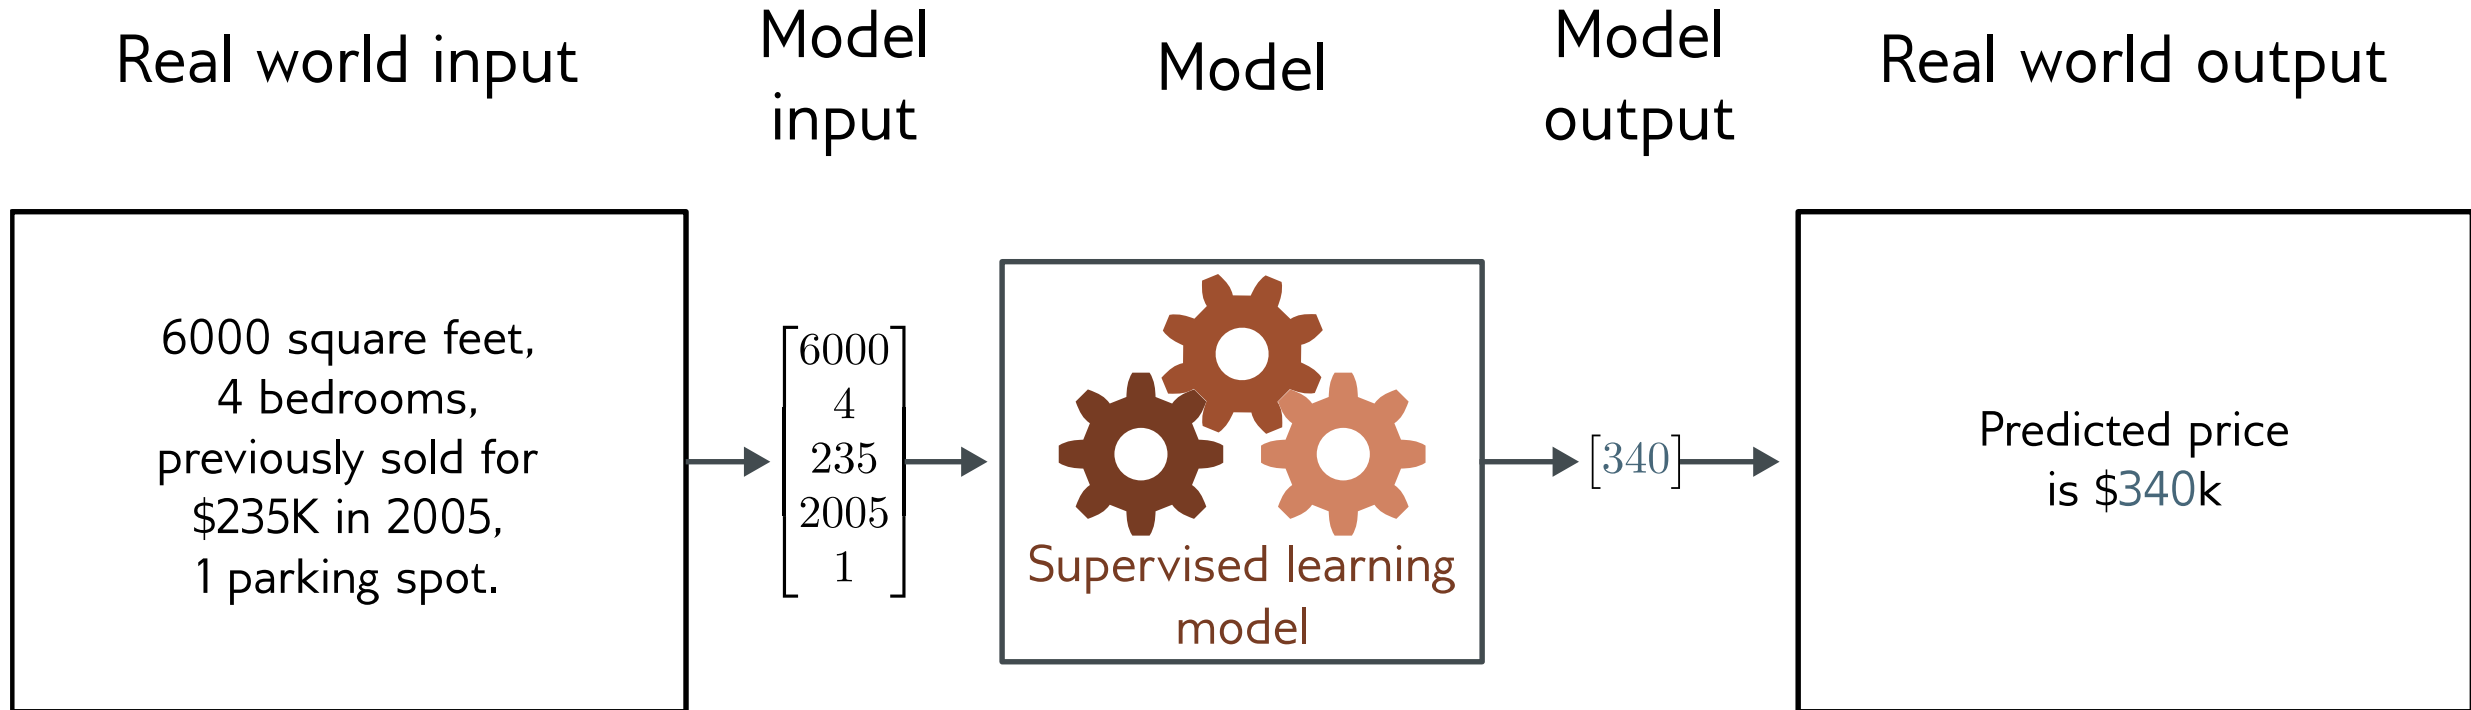


ML 101

- Approximate a function $y = f(x; \theta)$
 - x : image pixels; y : classes (dog, cat)
 - x : house size; y : house price
 - θ are the parameters
- Two typical problems
 - Classification: discrete categories
 - Regression: continuous categories

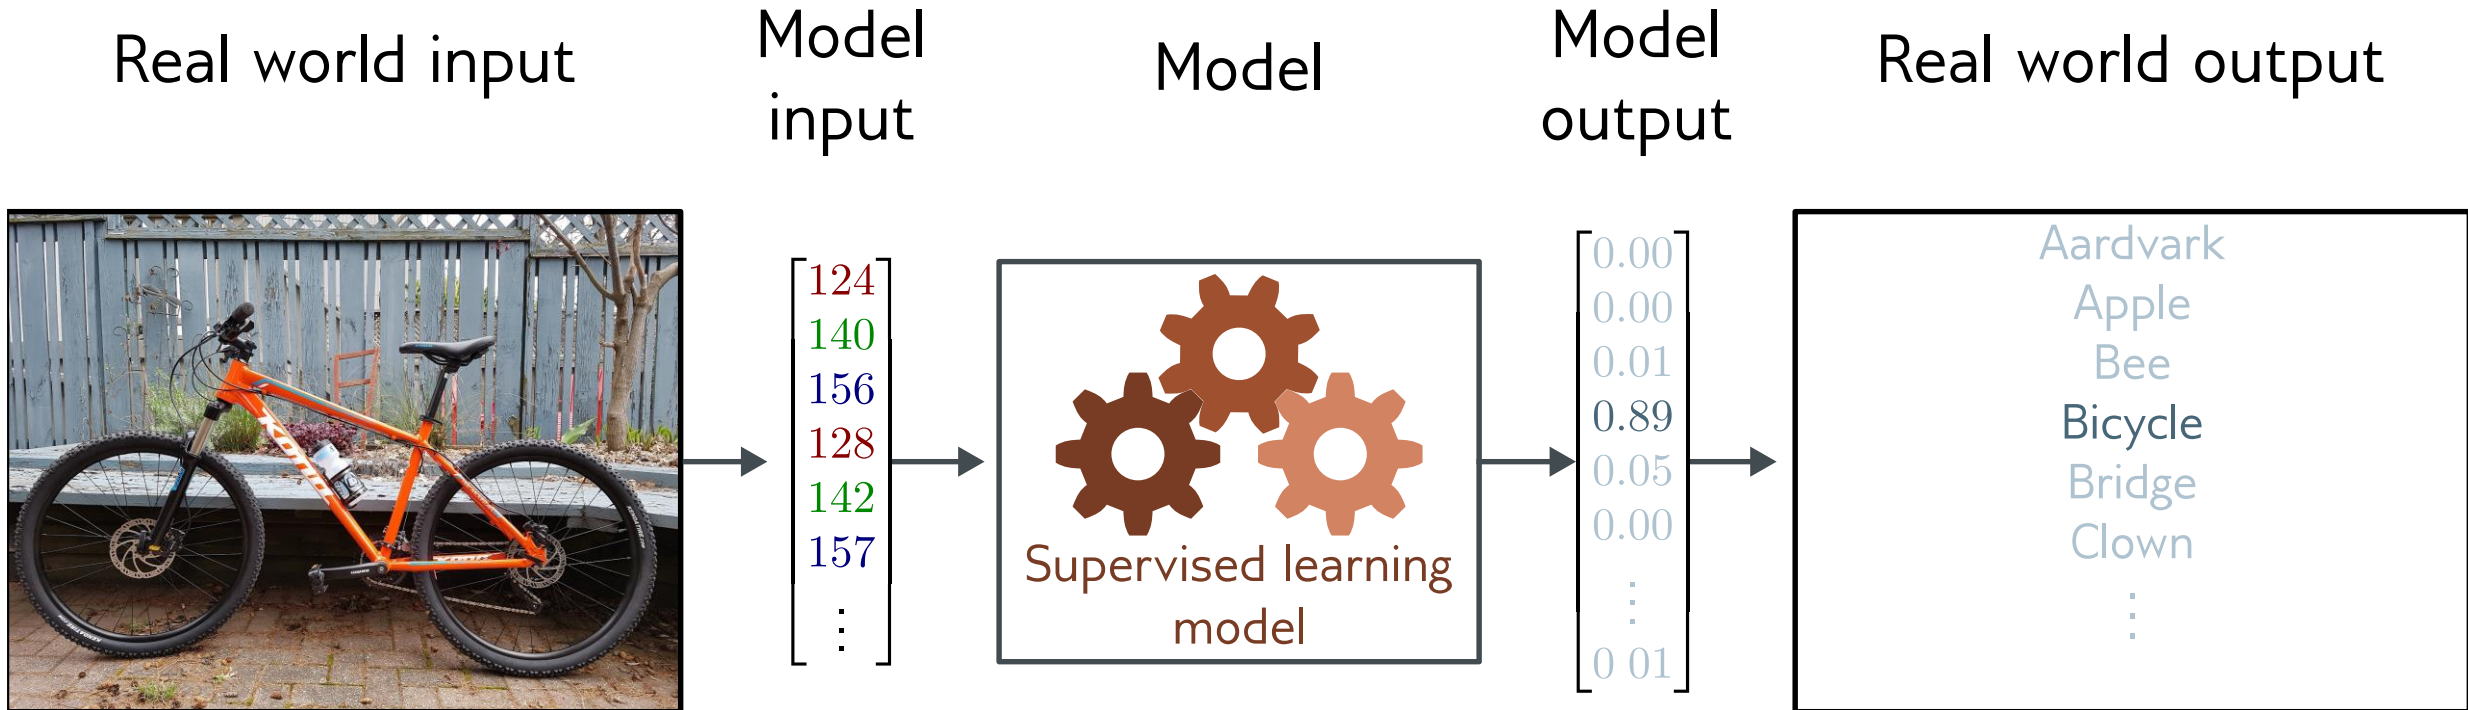


Regression



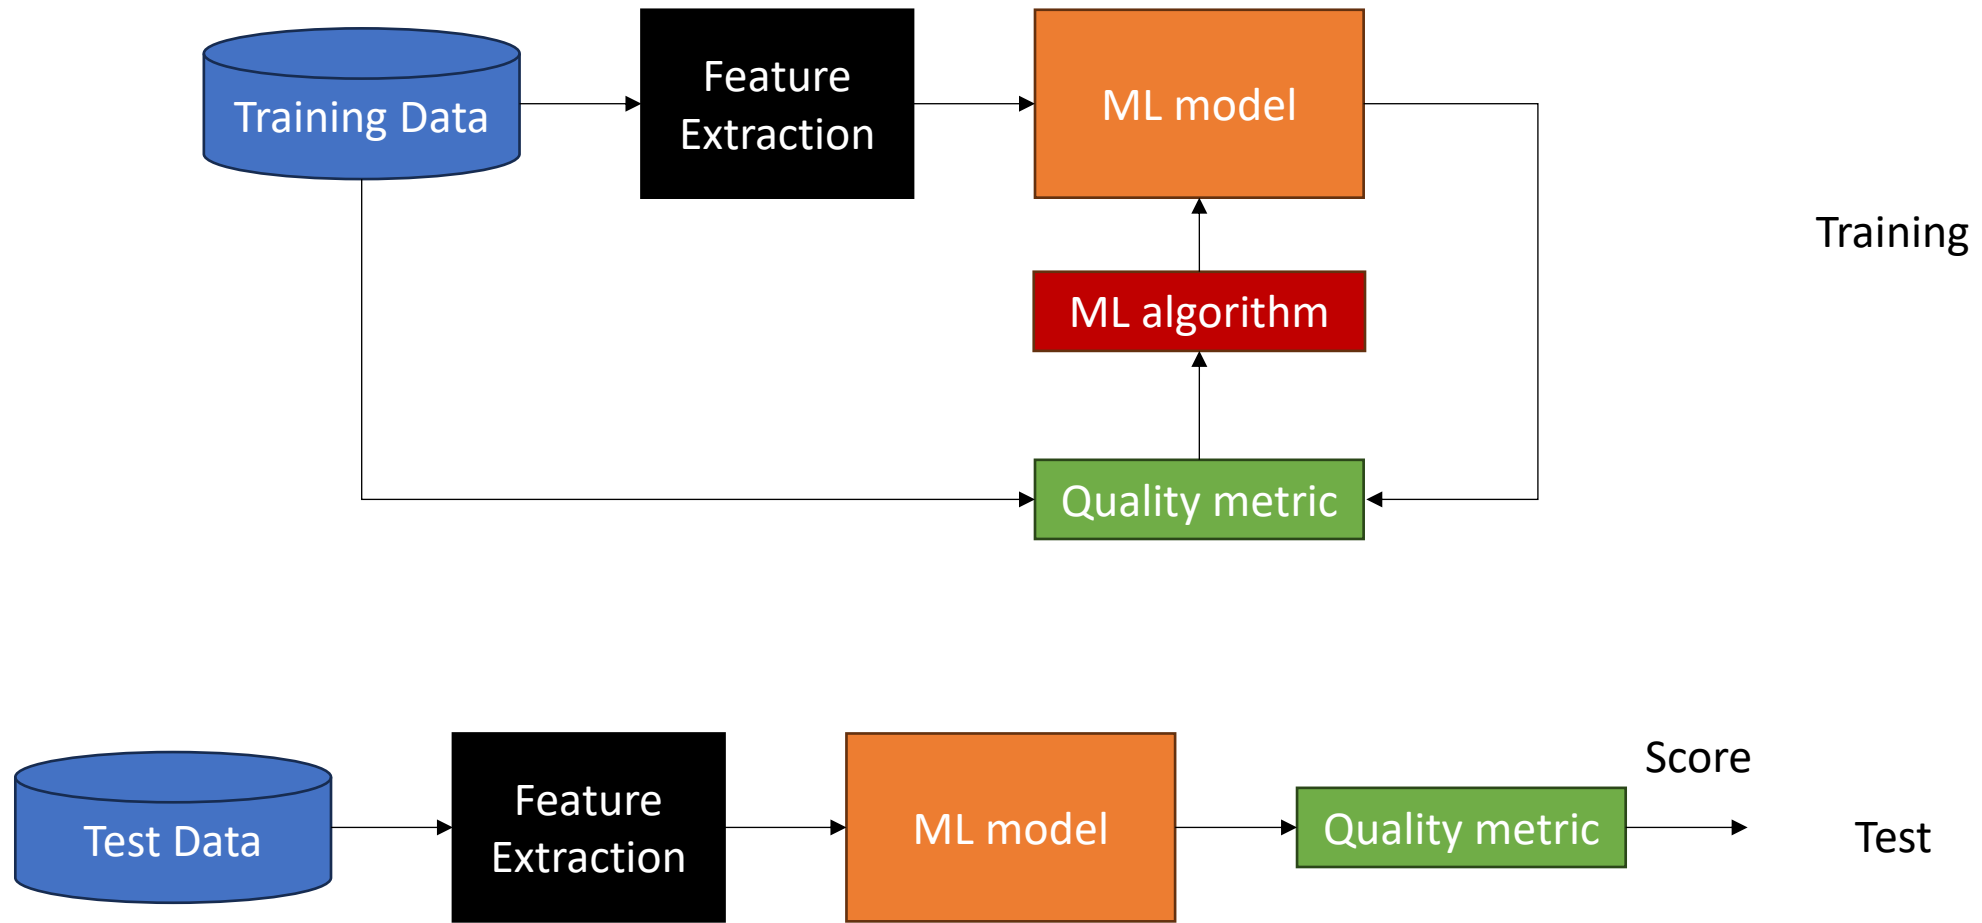
- Univariate regression problem (one output, real value)
- Fully connected network

Classification



- Multiclass classification problem (discrete classes, >2 possible classes)
- Convolutional network

ML workflow



<https://www.coursera.org/learn/ml-foundations?specialization=machine-learning>

Feature Extraction 101

- Machine learning needs numerical data as input:

$$\mathbf{x} \in \mathbb{R}^n$$

- Data normalization: assumption of data with mean 0, stddev 1

$$x' = \frac{x - \mu}{\sigma}$$

\$ x_std = StandardScaler().fit_transform(x)

- Typical representations

Data Type	Common Representation
Images	Pixel intensity arrays
Text	Word embeddings (Word2Vec, GloVe, BERT)
Audio	Spectrogram or MFCC features

Model Selection 101

- **Many ML models to choose from:**
 - **Linear models:** simple, interpretable, good baseline
 - **Decision Trees / XGBoost:** handle complex patterns, less linear assumptions
 - **Neural Networks:** powerful but need large data and tuning
- **Not every problem needs Deep Learning!**
 - NNs can overfit on small datasets
 - Many hyperparameters, hard to tune
 - Require more computation and time
- **Model selection = experimentation**
 - Try multiple models
 - Tune hyperparameters
 - Compare results using validation metrics (e.g., accuracy, RMSE, F1)
 - Pick the model that's **accurate, simple, and practical**



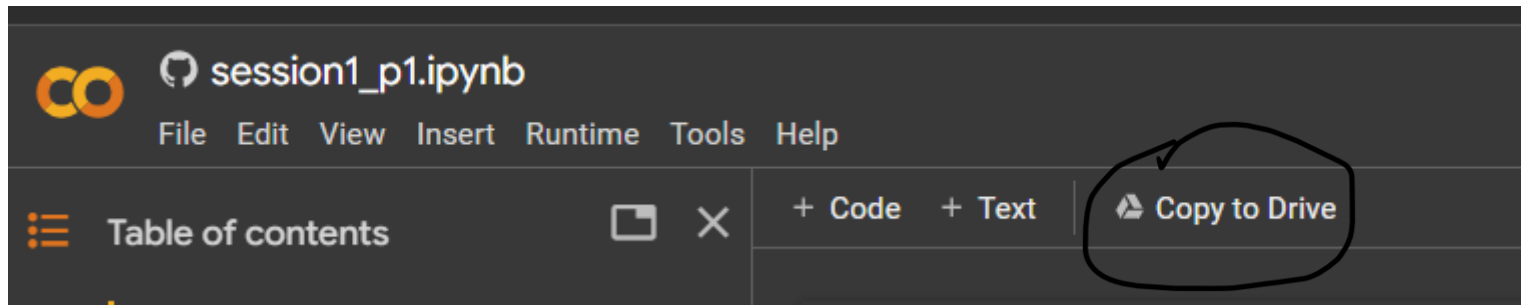
ChatGPT for everything XGBoost

Pro & Cons of DL

Pro	Cons
<ul style="list-style-type: none">• Can learn features & representations easily• Ability to process a lot of data• Flexible framework• Maps well into parallel hardware (GPU and others)	<ul style="list-style-type: none">• Hard to understand and build intuition of why the model works: explainability and interpretability• Requires a lot of data and expensive hardware.

Work session 1.0

Google Colab: Our tool for today

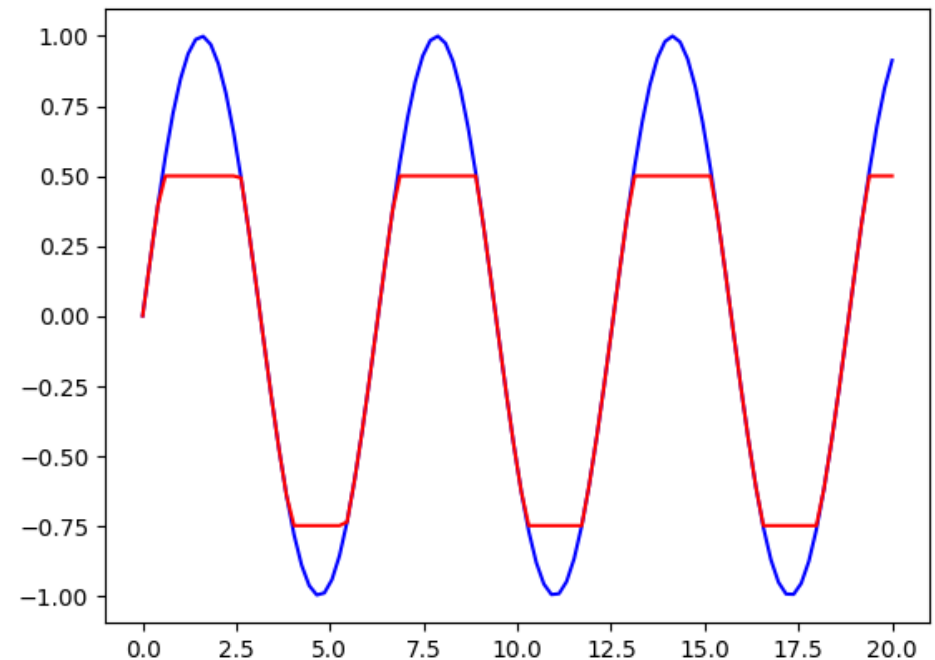


Handling numbers & data: numpy

- Numpy is library for handling array, vector & matrices
- Provide all keys operations in arrays
 - Creation : $a = np.array([[1, 2, 3], [4, 5, 6]])$
 - Add/subtract/multiply: $a + b$, $a - b$, $a * b$, ...
 - Dot product and matrix multiplication: $a @ b$
 - Slicing: $a[1] \Rightarrow [4, 5, 6]$

Plotting: matplotlib

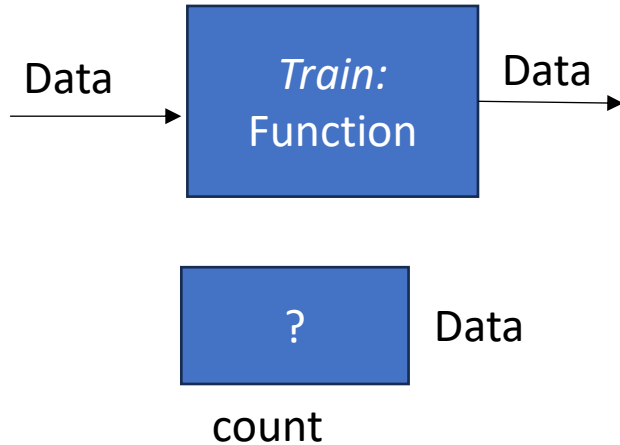
- Matplotlib is the Python library used to plot data
- You can create easily different types of plots (scatter plot, histogram, X-Y plot), add legend and different details.



Python 101: Why Python?

- A lot of libraries for ML/DL: PyTorch, scikit-learn, pandas, numpy
- Easy to learn, simple syntax
- Interactive notebooks: Jupyter Notebook, Kaggle, Google Collab
- Free & open-source

Python 101: Basic features



```
def vowel_count(word):  
    vowels = ["a", "e", "i", "o", "u"]  
    count = 0  
    for char in word: # loops  
        if char in vowels: # conditions  
            count += 1  
    return count
```

```
vowel_count("hello")
```

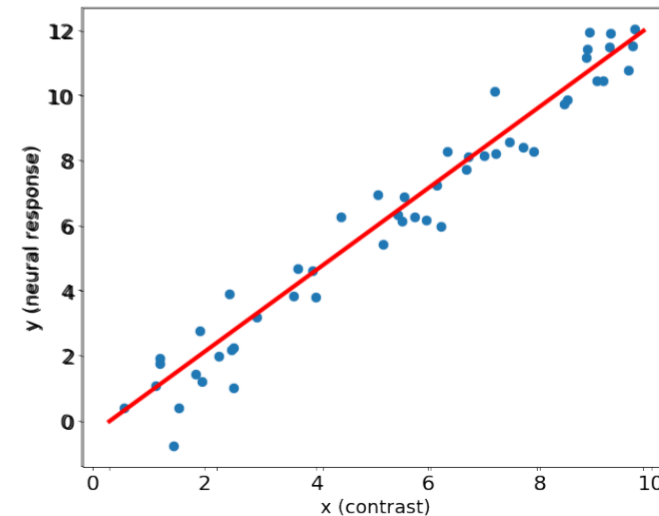
<i>MyModel: Class</i>
<i>my_model: Data1</i> <i>params: Data2</i> ...
<i>Train: Function1</i> <i>Test: Function2</i> ...

Linear Regression

Linear regression makes predictions about the linear relationship between the input variable x (contrast) and the output variable y (neural response).

$$\begin{array}{ccccccc} y & = & \theta_1 & \times & x & + & \theta_0 \\ \uparrow & & \uparrow & & \uparrow & & \uparrow \\ \text{neural response} & & \text{linear weight} & & \text{contrast} & & \text{Intercept} \end{array}$$

We are not considering the intercept for simplicity, resulting in a one-parameter model.



Linear Regression: MSE

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N (y_i - \theta x_i)^2$$

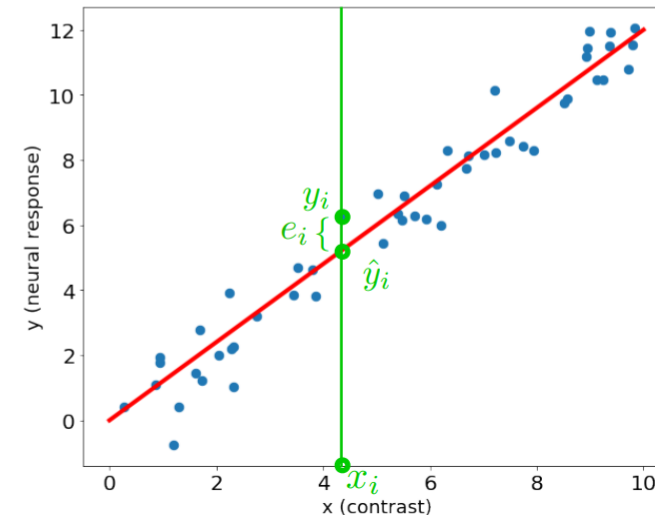
Mean Squared Error (MSE)

MSE computes the average error between the model prediction \hat{y} and the true y .

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 = \frac{1}{N} \sum_{i=1}^N e_i^2$$

Annotations for the MSE formula:

- $\frac{1}{N}$: total number of data points
- $\sum_{i=1}^N$: index of data points $i=1, \dots, N$
- y_i : true neural response
- \hat{y}_i : model prediction
- e_i : residual



Work session 2

How to Evaluate a Machine Learning Model (and Why It Matters)

Learning Objectives

- Preparing data (inspection, visualization, train/test splits)
- Building a machine learning model (Decision Tree)
- Choosing metrics to assess model performance
- Using the concept of bias–variance trade-off to tune hyperparameters
- **BONUS.** Reducing variance using ensembling

Link to Colab: <https://tinyurl.com/dlbs1p2>

Preparing Data

California Housing Dataset



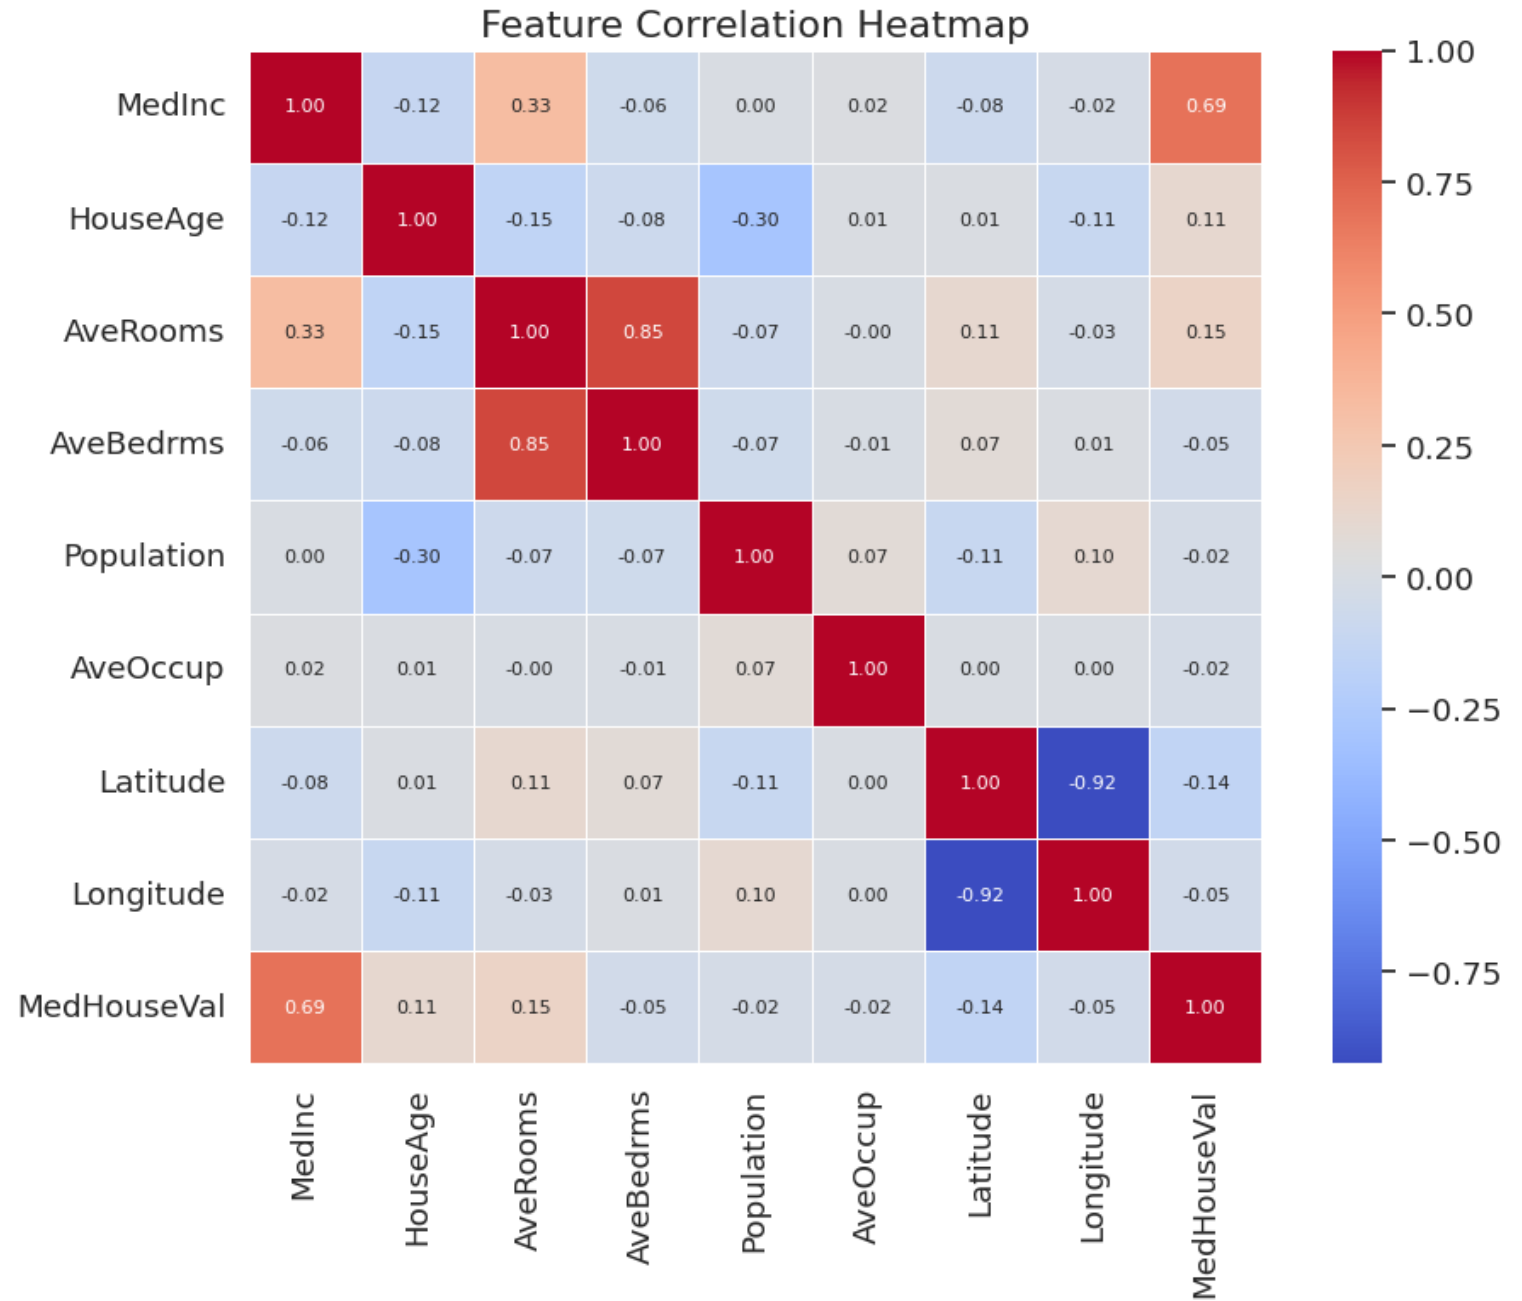
Preview of the dataset:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25

Activity 1

Q1. Which features are most correlated with the target?

Q2. Which features might be redundant or overlapping?



Transforming Data – *One-Hot Encoding*

Original data:

	id	species
0	1	cat
1	2	dog
2	3	snake
3	4	cat
4	5	dog
5	6	turtle
6	7	dog

One-hot encoded with `pandas.get_dummies`:

	id	is_cat	is_dog	is_snake	is_turtle
0	1	1	0	0	0
1	2	0	1	0	0
2	3	0	0	1	0
3	4	1	0	0	0
4	5	0	1	0	0
5	6	0	0	0	1
6	7	0	1	0	0

Transforming Data – *Feature Scaling*

```
Original (unscaled) data:
  temp_K  humidity  day_of_year
0   272.0     0.15         12
1   289.5     0.80        150
2   301.2     0.40        230
3   295.0     0.60        320
4   280.3     0.05         45
```

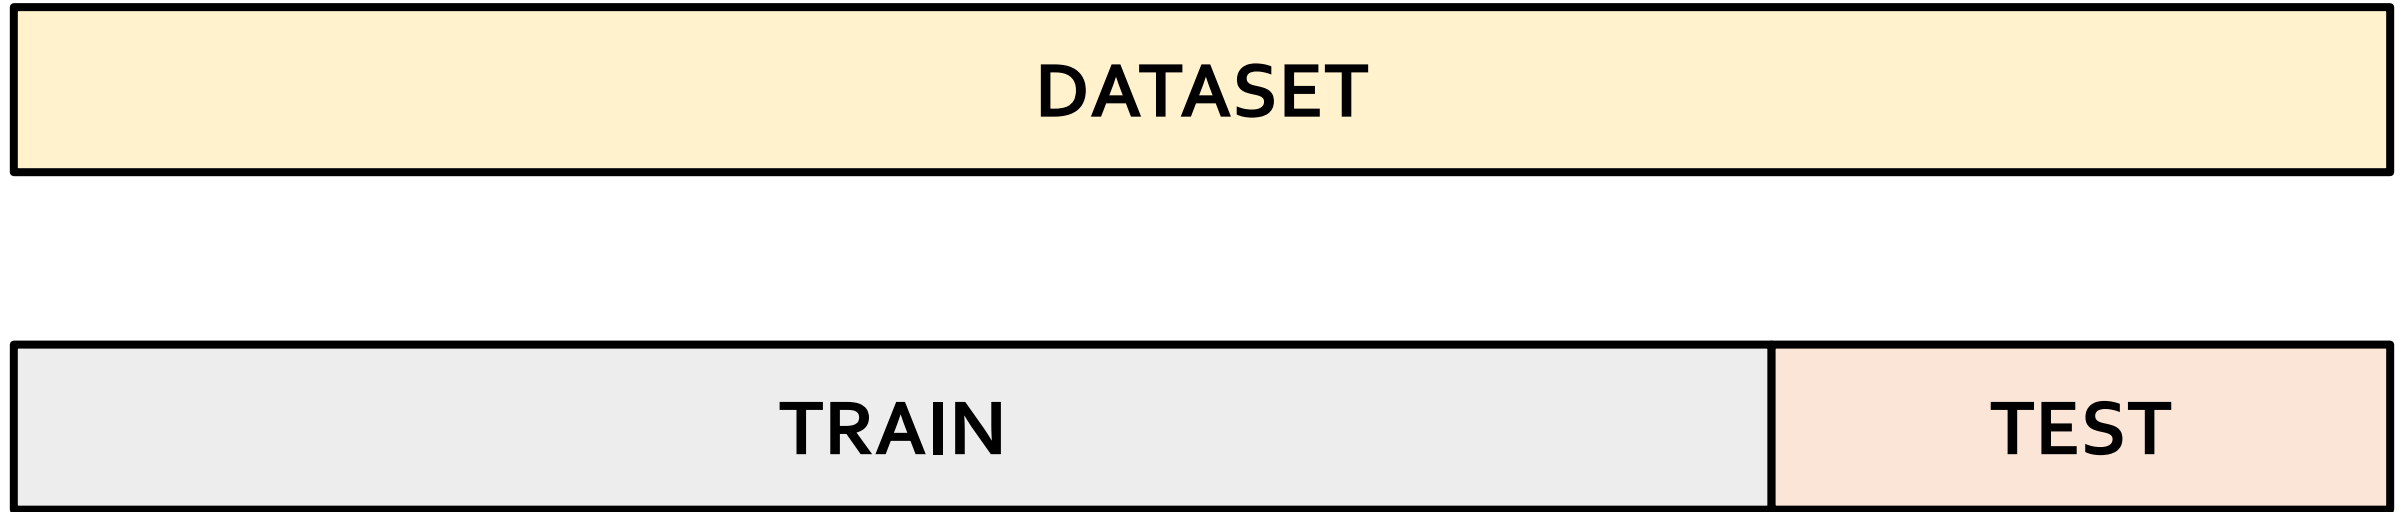
```
Standardized data (≈ mean 0, std 1):
  temp_K  humidity  day_of_year
0  -1.501    -0.901    -1.219
1   0.183     1.441    -0.012
2   1.309     0.000     0.687
3   0.712     0.721     1.475
4  -0.702    -1.261    -0.931
```

Standardization

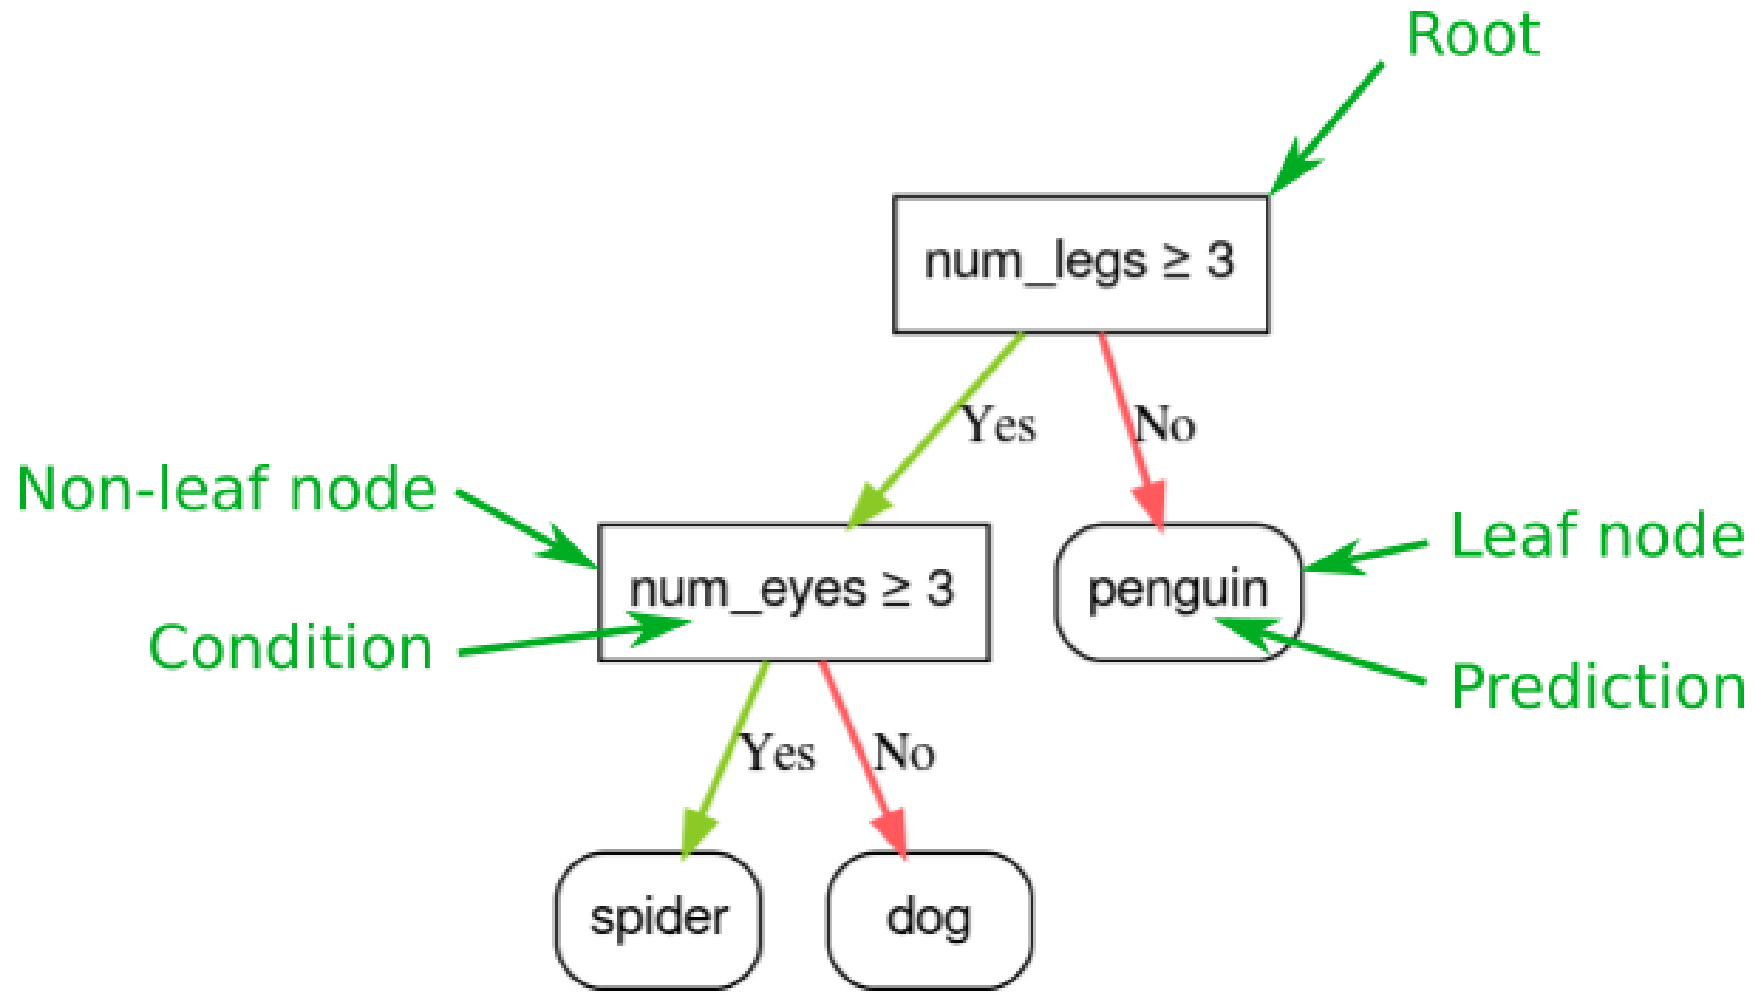
```
Min-Max scaled data (range ~[0, 1]):
  temp_K  humidity  day_of_year
0   0.000     0.133     0.000
1   0.599     1.000     0.448
2   1.000     0.467     0.708
3   0.788     0.733     1.000
4   0.284     0.000     0.107
```

Min-Max Scaling

Splitting Data for Training and Testing



Choosing a Machine Learning Model – *Decision Tree*



Model Evaluation Metrics – *Regression*

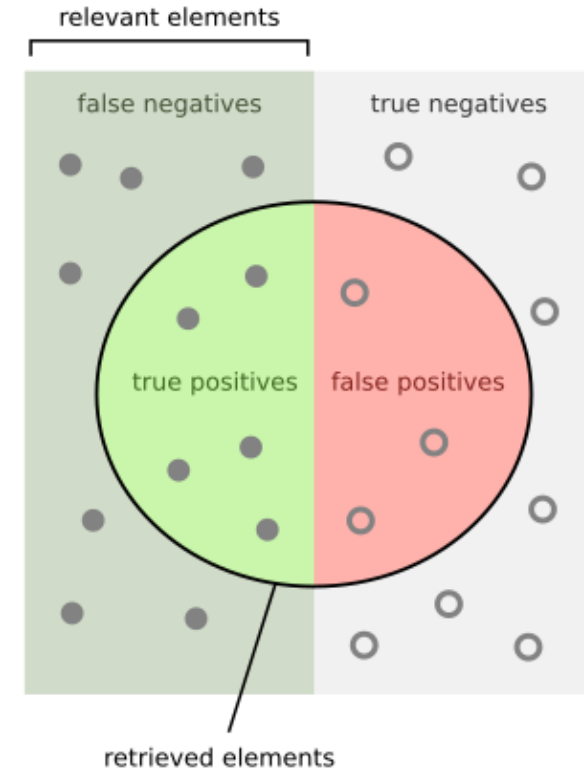
$$\text{RMSE}(y_{\text{true}}, y_{\text{pred}}) = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

$$R^2(y_{\text{true}}, y_{\text{pred}}) = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Activity 2. Write a function to compute RMSE for train and test splits

Model Evaluation Metrics – *Classification*

		Predicted condition	
		Cancer	Non-cancer
Actual condition	Total	8 + 4 = 12	
Cancer	8	6	2
Non-cancer	4	1	3



How many retrieved items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are retrieved?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Hyperparameter Tuning

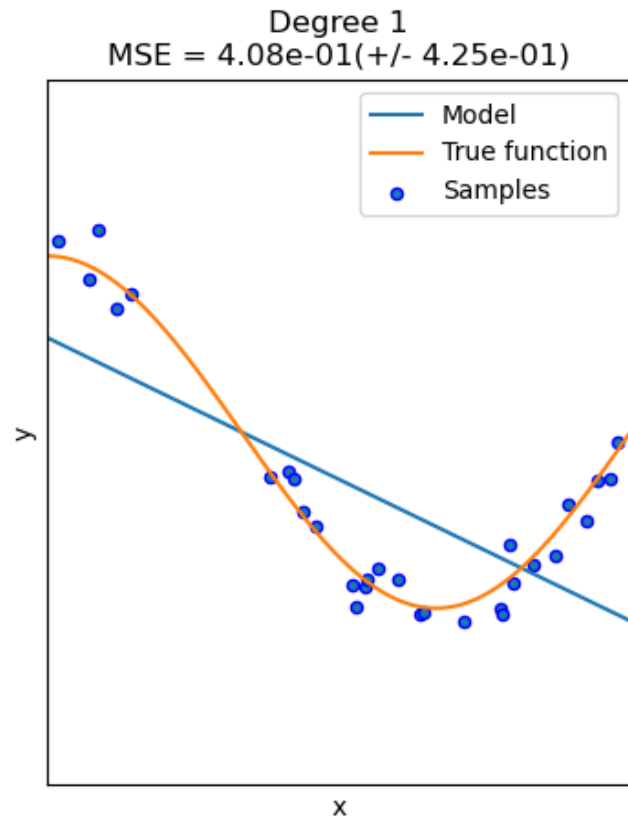
- **Hyperparameter.** A setting you choose before training that controls how the model learns. It's not *learned* from the data.
- Decision tree hyperparameters:
 - **max_depth** → limits how many “levels” the tree can grow
 - **min_samples_split** → minimum samples required to split a node
 - **min_samples_leaf** → minimum samples required in any leaf criterion

Activity 3. Which **max_depth** is the best?

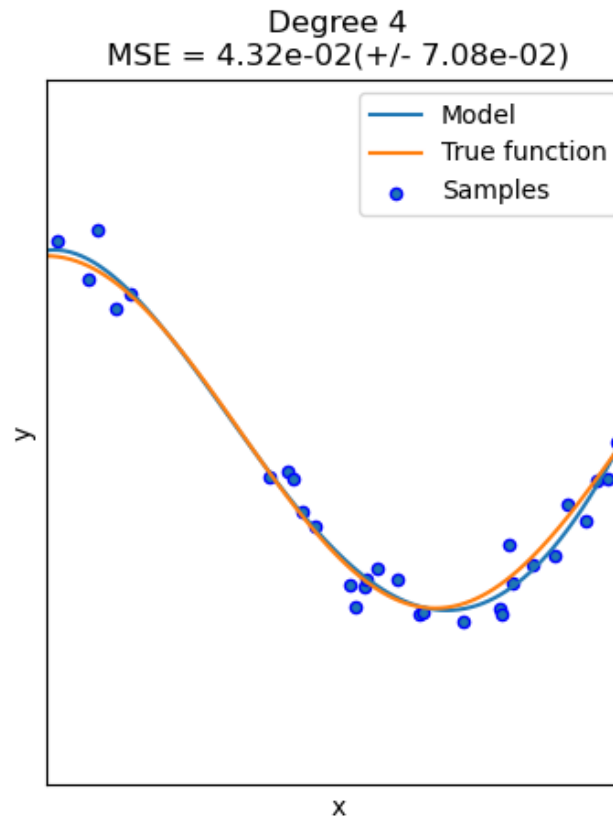
Splitting Data – *Enter Validation Split*



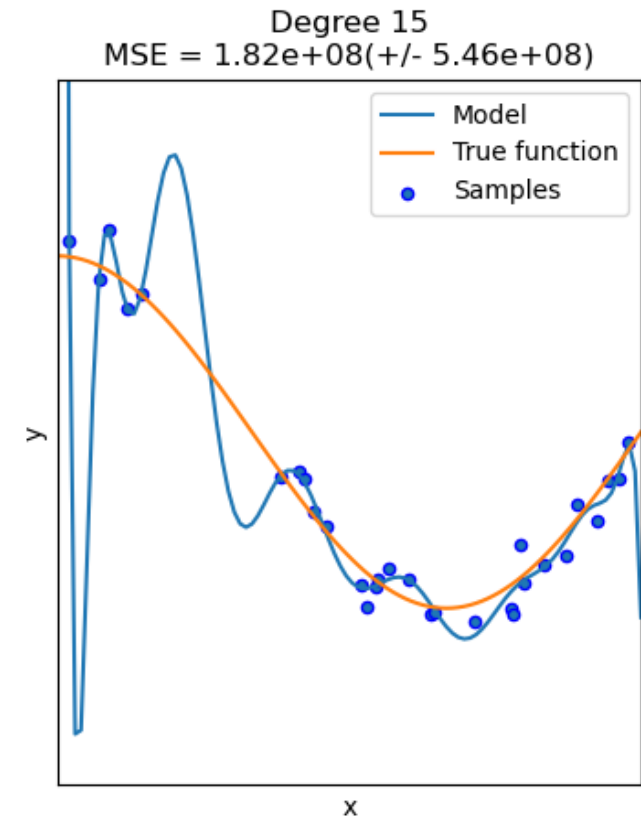
Bias-Variance Tradeoff



High bias
(Underfitting)

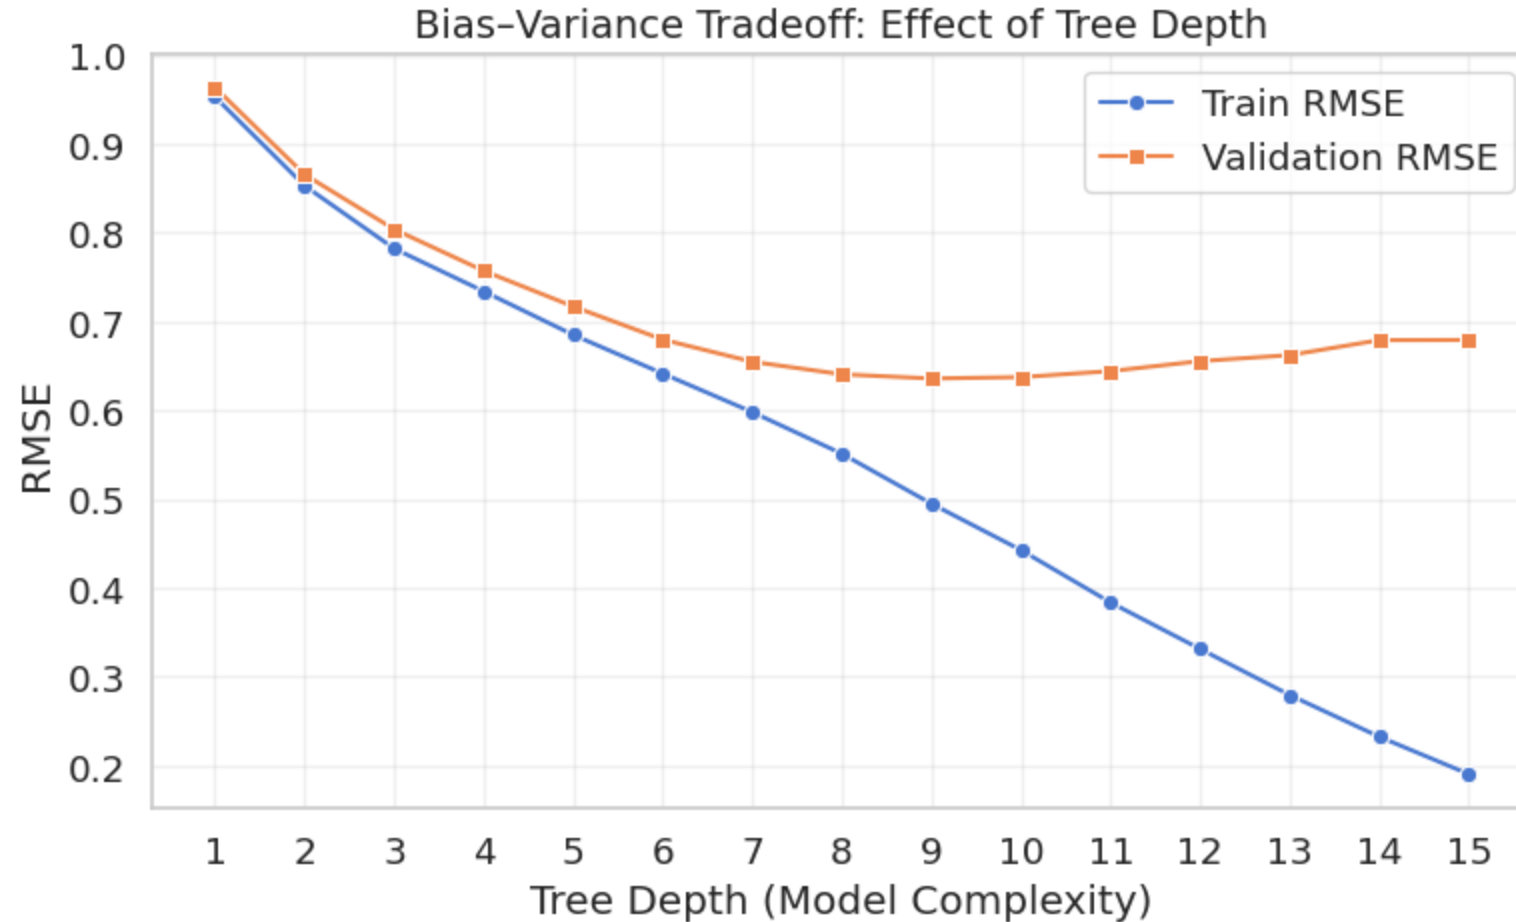


Good Fit



High Variance
(Overfitting)

Validation Curves



Activity 4. Which `max_depth` is the best?

Summary

Blueprint of a Machine Learning Project

1. Define your prediction goal (classification/regression)
2. Explore and understand your dataset
3. Split the data into train/test/validation
4. Pick a model, train it, and measure its performance
5. Tune hyperparameters and optimize model through the lens of bias–variance trade-off.

**We'd love to hear
your feedback**

Help us make the next
session even better!

**3rd DL Bootcamp (2025-26) -
Session 1 - Post Session Feedback**



References / Reading

- Python introduction: <https://swcarpentry.github.io/python-novice-inflammation/>
- More on scientific python: <https://lectures.scientific-python.org/>
- <https://deeplearning.neuromatch.io/>
- NeuroMatch Academy:
https://deeplearning.neuromatch.io/tutorials/W1D1_BasicsAndPytorch/chapter_title.html
- Exploratory computing w/ Python:
https://mbakker7.github.io/exploratory_computing_with_python/
- <https://udlbook.github.io/udlbook/>
 - Reuse slides from there

See you in
the next
session

DEEP LEARNING BOOTCAMP



DEEP LEARNING ESSENTIALS BOOTCAMP SESSION-2

Deep Learning for Regression and
Classification with PyTorch

02:30 PM, NOV 10TH, 2025

THE ARC TR+30

BRING YOUR LAPTOP !!!!



SRIRAM RANGA



SANTHISENAN AJITH

Upcoming Sessions	Date	Speakers
Deep Learning for Images	17 Nov 2025	Micheal Yuhas & Amashi Niwarthana
Natural Language Processing (NLP)	To be decided	To be decided



Scan Me

